Threshold Phases and the Limits of Motion Compensation for Event Streams

[Author list to be set]

***Abstract*—An event camera reports when the log intensity of a pixel crosses a threshold. When an edge moves across the sensor, every pixel on its path sees the same intensity profile, yet the pixels do not emit the same events: each fires on its own lattice of threshold levels, offset by a threshold phase, and every event carries timing jitter. We determine the rate that these differences require in an analytic model: a ramp edge that passes each pixel once under an affine flow, with independent uniform phases and Gaussian jitter. Given the edge and its motion, a lossless code must still describe the phase and the jitter of every pixel, so its rate per event does not fall as the path grows. Under a mean-squared timing tolerance the rate has three regimes. Below the jitter variance, the code describes every event time, and the rate is exact. Above it, the code describes one number per pixel, and the rate is known within 0.255 bit. Beyond a threshold set by the event spacing of the pixel, the rate given the edge and its motion is zero, and a description of the motion needs a number of nats that grows at most logarithmically with the length of the path. A sender has the events and not the edge, and the closest it can come is to move events along the motion. A moved event keeps the phase of its source pixel, so on a voxel grid the moved counts reproduce only a share of about Δ_t/τ of the true counts when the time bin Δ_t is shorter than the event spacing τ. A standard HEVC coder on count pictures follows these limits. On the model, its lossless P pictures cost about 0.1 bit per event when the bin is long against the event spacing and 5 to 14 bits when it is not. On recordings from a static camera and from a moving car, its lossless bits are 0.92 to 1.18 times those of a list of the counts without prediction, and on fine grids its quantizer removes single events instead of coarsening them.**

***Index Terms*—Event camera, send-on-delta sampling, level-crossing sampling, rate-distortion theory, generalized sampling, motion compensation, video coding, HEVC, neuromorphic sensing.**

I. Introduction

An event camera has no frames. Each pixel compares its log intensity with a stored reference and emits an event, with a time stamp and a polarity, whenever the difference reaches a contrast threshold [1], [2]. The output is a stream of address-events whose rate follows the scene: it is close to zero for a static scene and rises with the speed and the contrast of what moves. A sensor node that forwards this stream over a rate-limited link must decide how many bits the stream needs.

The stream of a moving scene looks highly redundant. If an object moves across the sensor, the pixels along the path of one of its edges observe the same intensity profile, one after the other. Yet they do not emit the same events. Each pixel compares the intensity with its own reference, so it fires on its own lattice of levels, and the offset of that lattice, the threshold phase of the pixel, differs from pixel to pixel. The same profile therefore produces a different event train at every pixel (Fig. 1). This paper asks what these differences cost: how many bits a code needs to preserve them, and how many it saves when a timing error is allowed.

The setting has precedents in sampling theory. An event pixel is a send-on-delta sampler [3], which for a monotone input coincides with level-crossing sampling [4], [5]. An object that crosses M pixels is sampled by M such samplers, which recalls the generalized sampling expansion [6], by which a bandlimited signal is recovered from the outputs of m linear systems that are each sampled at $1/m$ of the Nyquist rate. In multi-channel time encoding, several integrate-and-fire encoders with shifted integrators share one input, and the shifts are what make the channels complementary [7]–[9]. The threshold phases play that role here.

Existing event coders approach the rate empirically. Lossless coders exploit the spatial and temporal correlation of addresses and time stamps [10]–[12]. Lossy coders discard or aggregate events [13], [14], and the coder of [15] suppresses events that an on-sensor optical flow can predict and lets the receiver extrapolate them. A recent survey notes that reported results are hard to compare [16], and no accepted distortion measure for lossy event coding exists [17]. On the theoretical side, rate-distortion results for point processes concern Poisson and finite point processes [18]–[21]. They do not model events that a moving profile generates through a threshold rule. Event streams have also been aggregated into count pictures per polarity and coded without loss by HEVC [11], [22], a video coder whose prediction is motion compensated [23], but what the prediction saves on such pictures has not been reported.

This paper models that mechanism and derives what a code can and cannot save. The phase mechanism holds for any scene (Lemma 1). For the rates we specialize to a ramp edge that passes each pixel once under an affine flow, with independent uniform phases and Gaussian jitter, which isolates the mechanism and gives closed forms. A standard video coder then serves as a probe of the mechanism on recordings. The central result is the rate-distortion function of the events of a pixel under a timing tolerance. The other results lead to it or follow from it.

1) Mechanism. Given the edge and its motion, the events of a pixel are a function of a single number, its threshold phase, up to timing jitter. The phase is a property of the pixel and is independent across pixels. Next to what the pixels of a path share, the edge and its motion, each pixel therefore holds a private quantity. A lossless code must describe it for every pixel, and its rate per event does not decrease with the number M of pixels on the path (Theorem 1).

2) Rate under a timing tolerance. The tolerance decides which of the private quantities a code must describe (Theorem 2). Below the jitter variance it describes all of them,

and the rate is known exactly. Above it one number per pixel remains, and the rate is known within 0.255 bit. Beyond a threshold the rate given the edge and its motion is zero: the decoder regenerates the events. The threshold grows with the event spacing of the pixel, which is the contrast threshold divided by the rate of change of its intensity, so slowly crossed pixels need a larger tolerance. Along a path of M pixels, a description of the motion of $O(\ln M)$ nats then replaces a lossless description of order M.

3) Moved events and a video coder. A sender has the events and not the edge, and the closest it can come to regeneration is to move events along the motion, which is what the prediction of a video coder does. A moved event keeps the phase of its source pixel. At equal event spacings this doubles the timing error of the ideal decoder (Proposition 2), and on a voxel grid the moved counts reproduce a share of about Δ_t/τ of the true counts when the bin is shorter than the event spacing. On the model, the best block motion compensation follows this share within 0.005. We then code voxel counts of recordings from a static camera [24] and from a moving car [25] with x265, an unmodified HEVC encoder, and with a list of the counts that predicts nothing. Without loss, the motion-compensated prediction of x265 saves nothing against the list, and with loss its quantizer removes single events on fine grids.

Section II states the model. Sections III and IV contain the theory, Section V the moved events, and Section VI the measurements. Proofs are in the appendices. Rates are in nats unless bits are named, $\|\cdot\|$ is the Euclidean norm of a vector and the spectral norm of a matrix, and $[x]^+ = \max(x, 0)$.

II. Model

We first state the pixel rule, from which the phase mechanism follows for any scene. We then specialize to a ramp edge in uniform passage, for which the rates have closed forms.

A. Pixel

Pixels sit at the lattice points $i \in \mathbb{Z}^2$ of the image plane. Pixel i observes the log intensity $\ell_i(t) = L(i,t)$ of a continuous field L.

(A1) Send-on-delta. Pixel i holds a reference μ_i with $|\ell_i - \mu_i| < C$ at the start. It emits an event at the first time t at which $|\ell_i(t) - \mu_i| = C$, with polarity $p = \operatorname{sign}(\ell_i(t) - \mu_i)$, and then sets $\mu_i \leftarrow \mu_i + pC$. The contrast threshold $C > 0$ is common to all pixels and known to the decoder. The *phase* of the pixel is the number $a_i \in (0, C]$ with $\mu_i \in a_i + C\mathbb{Z}$ at the start.

B. Scene and Motion

A template $g : \mathbb{R}^2 \to \mathbb{R}$ gives the log intensity in object coordinates. At time t the object point u is at the image position $\phi_t(u) = J_t u + b_t$, with $J_t \in \mathbb{R}^{2\times 2}$ invertible, so

$$L(x,t) = g\big(\phi_t^{-1}(x)\big) \tag{1}$$

and pixel i observes

$$\ell_i(t) = g\big(u_i(t)\big), \qquad u_i(t) = J_t^{-1}(i - b_t). \tag{2}$$

The curve u_i is the trajectory of pixel i in object coordinates. The velocity field $v(x,t) = \dot\phi_t(\phi_t^{-1}(x))$ is affine in x and satisfies the brightness constancy relation

$$\frac{\partial L}{\partial t}(x,t) + \langle \nabla L(x,t), v(x,t)\rangle = 0, \tag{3}$$

which follows from differentiating $L(\phi_t(u), t) = g(u)$ in t. We call the motion an *affine flow* when the velocity field is also constant in time,

$$v(x) = \Lambda x + v_0, \tag{4}$$

with generator $\Lambda \in \mathbb{R}^{2\times 2}$. A translation has $\Lambda = 0$. The trace of Λ is the rate at which the object grows or shrinks in area, its antisymmetric part rotates the object, and its symmetric traceless part shears it. At a given instant, any smooth velocity field has the form (4) to first order in x around a point, so (4) is the local model of a general motion over a short time.

C. Edge

(A2) Ramp edge. The template is a ramp across a line with unit normal ν,

$$g(u) = \varrho\big(\langle \nu, u\rangle - c_0\big), \qquad \varrho(z) = \begin{cases} 0, & z \le 0, \\ \gamma z, & 0 \le z \le w, \\ V, & z \ge w, \end{cases} \tag{5}$$

with slope $\gamma > 0$, width w and height $V = \gamma w = nC$ for an integer $n \ge 1$.

Pixel i sees the edge through its edge coordinate

$$z_i(t) = \langle \nu, u_i(t)\rangle - c_0, \qquad \ell_i(t) = \varrho\big(z_i(t)\big). \tag{6}$$

(A2′) Uniform single passage. The edge passes pixel i once: $z_i < 0$ before a time t_i, $z_i > w$ after the passage, and in between z_i is affine in t with slope $r_i > 0$.

Inside the ramp, $\nabla L = \gamma \nu_t$ with $\nu_t = J_t^{-\top}\nu$ a normal of the edge in the image, and $\dot z_i = -\langle \nu_t, v(i,t)\rangle$. By (3), the intensity of pixel i rises during the passage at the rate

$$\gamma r_i = -\langle \nabla L, v\rangle = \|\nabla L\|\, v_{\perp,i}, \tag{7}$$

with the field evaluated at pixel i. Here $v_{\perp,i} = -\langle \nabla L, v\rangle/\|\nabla L\|$ is the speed at which the edge advances over the pixel along its normal, counted positive when the intensity of the pixel rises, as it does in a passage with $r_i > 0$. We call

$$\tau_i = \frac{C}{\gamma r_i} = \frac{C}{\|\nabla L\|\, v_{\perp,i}} \tag{8}$$

the *event spacing* of pixel i. It is the time that the ramp takes to pass the pixel, divided by n. At a fixed image slope $\|\nabla L\|$ it is inversely proportional to the normal speed of the edge. Assumption (A2′) holds exactly for a translation, where τ_i is the same for all pixels, and for a plane that approaches the camera along the optical axis at constant speed, where J_t^{-1} is affine in t and b_t is fixed. For an affine flow it holds to first order (Remark 1), and τ_i varies from pixel to pixel.

D. Phase, Jitter and Clock

(A3) Phase. Before the edge arrives, pixel i rests at the level 0 with its reference behind the passage, $\mu_i \in (-C, 0]$, so that $a_i = \mu_i + C$. The phases a_i are independent across pixels, uniform on $(0, C]$, and independent of the jitter.

(A4) Jitter and clock. The kth event of pixel i is recorded at $T_{i,k} = t_{i,k} + \varepsilon_{i,k}$, where $t_{i,k}$ is the time given by (A1) and the $\varepsilon_{i,k}$ are independent $\mathcal{N}(0, \sigma^2)$ variables, independent of the phases, with $\sigma > 0$. The stored value is $Y_{i,k} = \delta\lfloor T_{i,k}/\delta \rfloor$ for a clock resolution δ. The index k of each recorded event is known.

The uniform phase does not arise inside the model, in which every pixel rests at the same level. It stands for the offsets of real pixels and for their earlier activity. A pixel that the edge crosses from the high level to the low level is covered by the mirrored assumption, with its reference in $[V, V + C)$. Remark 4 treats a reference that is ahead of the passage.

We write $T_i = (T_{i,1}, \dots, T_{i,n})$ and $Y_i = (Y_{i,1}, \dots, Y_{i,n})$, and $\mathcal{P}$ for the set of pixels that the edge passes during the observation, with $M = |\mathcal{P}|$. Let $\eta(a, \sigma)$ denote the differential entropy of the sum of a uniform variable on $(0, a)$ and an independent $\mathcal{N}(0, \sigma^2)$ variable. It satisfies

$$\tfrac{1}{2}\ln\big(a^2 + 2\pi e\sigma^2\big) \le \eta(a, \sigma) \le \tfrac{1}{2}\ln\big(2\pi e(a^2/12 + \sigma^2)\big) \quad (9)$$

by the entropy power inequality [26, Ch. 17] and by the Gaussian bound on the entropy at a given variance [26, Ch. 8], and $\eta(a, \sigma) - \ln a \to 0$ as $\sigma/a \to 0$.

III. Structure and Lossless Limit

Lemma 1. *Under (A1), for any continuous ℓ_i, every event of pixel i occurs at a time t with $\ell_i(t) \in a_i + C\mathbb{Z}$.*

The lattice $a_i + C\mathbb{Z}$ never changes. The phase is a property of the pixel, set by its history and not by the edge that passes.

Lemma 2. *Under (A1), (A2), (A2′) and $\mu_i \in (-C, 0]$ before the passage, pixel i emits exactly n events, all of positive polarity, at*

$$t_{i,k} = t_i + A_i + (k-1)\tau_i, \qquad k = 1, \dots, n, \quad (10)$$

where $A_i = a_i/(\gamma r_i)$. Under (A3), A_i is uniform on $(0, \tau_i]$.

Given the template and the flow, the noise-free events of a pixel are thus a function of one number, its phase. Fig. 1 shows two pixels with $n = 3$ that see the same ramp: the intensity trajectories coincide, and the event trains differ by the difference of the phases. In the object coordinate, an event is a point of the graph of the profile on the level lattice of its pixel: the M pixels of a path sample the same profile on M lattices with independent offsets.

Theorem 1 (Lossless code length)**.** *Under (A1) to (A4) and (A2′), the vectors T_i, $i \in \mathcal{P}$, are independent given the template and the flow,*

$$h(T_i) = \frac{n-1}{2}\ln(2\pi e\sigma^2) + \eta(\sqrt{n}\,\tau_i, \sigma), \quad (11)$$

and

$$H(Y_i \mid g, \phi) = n\ln\frac{1}{\delta} + h(T_i) + e_\delta, \quad (12)$$

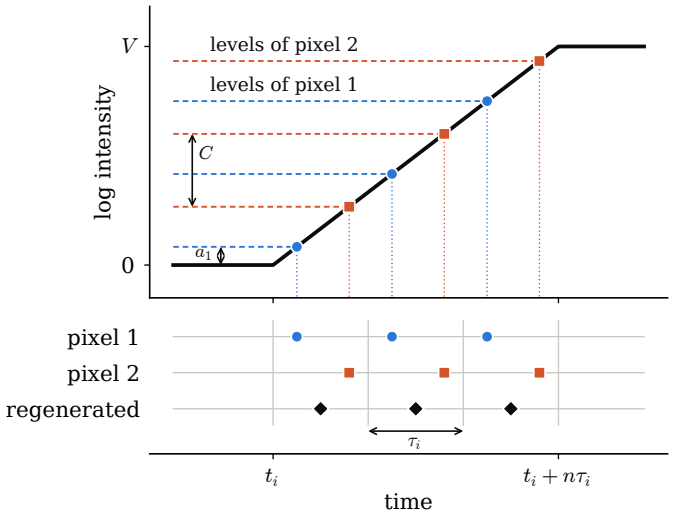


Fig. 1. One ramp, two pixels with different threshold phases, and the regenerated events (22). Time is measured from the arrival of the edge at each pixel. Jitter is omitted.

where $e_\delta \ge 0$ for every δ, and $e_\delta \to 0$ as $\delta/\sigma \to 0$, uniformly in τ_i.

When $\sigma \ll \tau_i$ and $\delta \ll \sigma$, (12) gives the cost per event

$$\frac{H(Y_i \mid g, \phi)}{n} \approx \frac{n-1}{n}\ln\frac{\sigma\sqrt{2\pi e}}{\delta} + \frac{1}{n}\ln\frac{\sqrt{n}\,\tau_i}{\delta}. \quad (13)$$

One event of each pixel carries the cost of the phase, and the other $n-1$ carry that of the jitter. The code length of the path is the sum of (12) over $\mathcal{P}$. It grows in proportion to M, and neither term of (13) falls as the path gets longer: a decoder that knows the edge and its motion exactly still needs (13) for every event. What the pixels of a path share, the edge and its motion, a code can describe once. What is private to each pixel, its phase and its jitter, it must describe for each pixel, and (12) is that irreducible conditional rate. Section VI-E measures what the motion saves a video coder on recordings.

Remark 1 (Accuracy of the uniform passage). For an affine flow (4), $\dot{\nu}_t = -\Lambda^\top \nu_t$ and $\ddot{z}_i = \langle \nu_t, \Lambda v(i)\rangle$, hence $|\ddot{z}_i| \le \kappa_i |\dot{z}_i|$ with $\kappa_i = \|\Lambda\|/|\cos\vartheta_i|$, where ϑ_i is the angle between the velocity $v(i)$ and the normal ν_t. Take r_i as the slope of z_i at t_i. Over a passage of duration $P_i = n\tau_i$, the event times then deviate from (10) by at most $\kappa_i P_i^2/2$, to first order in $\kappa_i P_i$. For a generator of norm $0.3\,\mathrm{s}^{-1}$, a passage of 5 ms and motion along the normal this is $4\,\mu$s, three orders of magnitude below the passage time. The bound grows without limit for edges that are nearly parallel to their motion. Such edges pass few pixels per unit time and contribute few events.

Remark 2 (Several edges). Theorem 1 counts one edge that passes each pixel once. If further edges of the same object pass the same pixel, they meet the same lattice, so the phase is paid once per pixel and not once per edge. With many edges per pixel the cost per event tends to the jitter term $\ln(\sigma\sqrt{2\pi e}/\delta)$, which does not depend on the path either.

IV. Rate Under a Timing Tolerance

The results of this section and the next assume decoders with different inputs. Theorems 2 and 3 count the bits of the event times when the encoder and the decoder know the template, the flow, and hence the pixels that the edge passes and the number of events of each, and when the index of every event is known (A4). Their zero rate is a rate given all of this. Proposition 1 adds the cost of the flow, while the decoder still holds the template. A practical sender holds neither, and Section V turns to what it can send instead.

The decoder outputs $\hat{T}_i \in \mathbb{R}^n$ for each pixel, and the distortion is the mean-squared timing error

$$d(T_i, \hat{T}_i) = \frac{1}{n}\|T_i - \hat{T}_i\|^2. \tag{14}$$

A. One Pixel

Let $R_i(D) = \inf I(T_i; \hat{T}_i)$, with the infimum over the conditional laws of $\hat{T}_i$ that satisfy $\mathbb{E}\, d(T_i, \hat{T}_i) \le D$, be the rate-distortion function of the events of pixel i. It is the rate that block codes over many independent passages approach. Define

$$s_i^2 = \frac{n\tau_i^2}{12} + \sigma^2, \qquad D_{A,i} = \sigma^2 + \frac{\tau_i^2}{12}, \tag{15}$$

$$\Gamma_i = \tfrac{1}{2}\ln(2\pi e\, s_i^2) - \eta(\sqrt{n}\,\tau_i, \sigma). \tag{16}$$

By (9), $0 \le \Gamma_i \le \frac{1}{2}\ln(2\pi e/12)$, which is 0.255 bit, for every σ and τ_i.

Theorem 2 (Rate-distortion function of a pixel)**.** *Under (A1) to (A4) and (A2′),*

$$R_i(D) = h(T_i) - \frac{n}{2}\ln(2\pi e D) \qquad \textit{for } 0 < D \le \sigma^2, \tag{17}$$

$$\big[U_i(D) - \Gamma_i\big]^+ \le R_i(D) \le U_i(D) \qquad \textit{for } \sigma^2 < D < D_{A,i}, \tag{18}$$

$$R_i(D) = 0 \qquad \textit{for } D \ge D_{A,i}, \tag{19}$$

with

$$U_i(D) = \frac{1}{2}\ln\frac{s_i^2}{nD - (n-1)\sigma^2}. \tag{20}$$

Moreover,

$$R_i(D) \ge \frac{n\,(D_{A,i} - D)}{2s_i^2} \qquad \textit{for } 0 < D \le D_{A,i}, \tag{21}$$

so that $R_i(D) > 0$ for $D < D_{A,i}$.

The three regimes have a direct reading. Below the jitter variance, every bit of timing precision costs a bit, per event: the rate falls by $\log_2 10 = 3.32$ bits per event and per decade of the root-mean-square tolerance $\sqrt{D}$. Between σ^2 and $D_{A,i}$, the $n-1$ jitter components orthogonal to a common shift of the events are no longer described, and the code describes one number per pixel, the phase together with the mean jitter, to within 0.255 bit. From $D_{A,i}$ on, nothing is described. The decoder that attains (19) outputs the mean times

$$\hat{T}_{i,k} = t_i + \big(k - \tfrac{1}{2}\big)\tau_i, \qquad k = 1, \dots, n. \tag{22}$$

It *regenerates* the events of the pixel from the edge and the flow. Each regenerated event is within $\tau_i/2$ of its noise-free time, and the count and the polarities are exact.

Two of these statements extend beyond the ramp. If the noise-free event times of a pixel are any function of its phase, (17) holds below the jitter variance with the entropy of the actual event times, and the smallest distortion at zero rate is $\frac{1}{n}\operatorname{tr}\operatorname{Cov}(T_i)$, attained by the mean times. The middle regime, in which one number per pixel is described to within 0.255 bit, and the closed forms (11), (20) and (15) belong to the ramp in uniform passage, where the phase shifts all events of a pixel by a common time.

The threshold $D_{A,i}$ grows with the square of the event spacing. By (8), a pixel whose intensity changes slowly has a large spacing and needs a large tolerance before its events can be regenerated, and a pixel that is crossed quickly needs a small one.

B. A Path

For the M pixels of a path, let the distortion be the average of (14) over $i \in \mathcal{P}$, and let $R_{\mathcal{P}}(D)$ be the rate-distortion function of $(T_i)_{i\in\mathcal{P}}$ under it.

Theorem 3 (Rate-distortion function of a path)**.** *Under (A1) to (A4) and (A2′),*

$$R_{\mathcal{P}}(D) = \min\Big\{ \sum_{i\in\mathcal{P}} R_i(D_i) : \ \frac{1}{M}\sum_{i\in\mathcal{P}} D_i \le D \Big\}. \tag{23}$$

In particular,

$$R_{\mathcal{P}}(D) = \sum_{i\in\mathcal{P}} h(T_i) - \frac{Mn}{2}\ln(2\pi e D) \qquad \textit{for } 0 < D \le \sigma^2, \tag{24}$$

and $R_{\mathcal{P}}(D) = 0$ if and only if

$$D \ \ge \ \bar{D}_A = \sigma^2 + \frac{1}{12M}\sum_{i\in\mathcal{P}} \tau_i^2. \tag{25}$$

If $\tau_i \le \tau_{\max}$ on the path, then for $D \le \bar{D}_A$

$$R_{\mathcal{P}}(D) \ge \frac{nM\,(\bar{D}_A - D)}{2\,(n\tau_{\max}^2/12 + \sigma^2)}. \tag{26}$$

If the tolerance is imposed on every pixel separately, the rate is zero if and only if $D \ge \max_i D_{A,i}$, and the slowest crossing sets the tolerance. By (26), below $\bar{D}_A$ a rate that is proportional to M remains.

C. The Motion

Theorems 2 and 3 assume that the decoder knows the flow. It remains to count what describing the flow costs. An error in a speed displaces the regenerated events in proportion to the distance that the edge has traveled, so a longer path needs the motion to a finer precision for the same distortion. A parameter of fixed dimension costs only logarithmically more at a finer precision, and Proposition 1 states this in general form.

Proposition 1 (Description of the motion)**.** *Let the flow depend on a parameter $\theta = (\theta_1, \dots, \theta_q)$ that ranges over a box whose jth side has length at most $L_j M^{\alpha_j}$, and let the mean times*

$\bar{t}_{i,k}(\theta) = t_i(\theta) + (k - \frac{1}{2})\tau_i(\theta)$ *of* (22) *satisfy* $|\partial \bar{t}_{i,k}/\partial\theta_j| \le S_j M^{\beta_j}$ *on the box, for all events of the path, with constants* $L_j, S_j > 0$ *and exponents* $\alpha_j, \beta_j \ge 0$ *that do not depend on* M*. Then for every* $\xi > 0$ *the events of the path can be described with average distortion at most* $(1+\xi)\bar{D}_A$ *using*

$$B \le \Big(\sum_{j=1}^{q} (\alpha_j + \beta_j) \Big) \ln M + O(1) \quad \textit{nats,} \tag{27}$$

where the $O(1)$ *term depends on q,* ξ*,* σ *and the constants and not on* M*.*

For a translation along a line with a speed in a fixed interval that excludes zero, θ holds the position of the edge, with a range of order M and a sensitivity of order one, and the speed, with a range of order one and a sensitivity of order M, so $B \le 2\ln M + O(1)$. The same holds with finite exponents, and hence $B = O(\ln M)$, for an affine flow whose normal speed is bounded away from zero on the path and whose duration is polynomial in M. The lossless cost of the same events is the sum of (12), of order M, so the ratio of the two grows at least as $M/\ln M$. Above $\bar{D}_A$ the description that remains is $O(\ln M)$. Below it, a rate proportional to M remains. The bound (27) is for a receiver that holds the edge.

Remark 3 (Event count under a change of scale). An edge segment emits n events for every pixel that it passes, so its event rate is n times the area that it sweeps per unit time, that is n times its length times its mean normal speed. Under a similarity flow about a fixed point, growth by a factor ζ makes the segment ζ times longer and ζ times faster, and the rate grows as ζ^2. For an object that approaches at constant speed, the image speed itself grows as ζ^2 and the rate as ζ^3. A regeneration that moves a fixed set of events keeps their number. Under a change of scale it must also change the number of events, or regenerate them from the edge.

Remark 4 (Departures from the model). *Reference ahead of the passage.* If $\mu_i \in (0, C)$ before a rising passage, as after a passage of the opposite polarity at the same pixel, the first level that the intensity meets is the reference itself, which does not fire. The pixel then emits $n-1$ events, and the results hold with $n-1$ in place of n and $t_i + \tau_i$ in place of t_i. *Height.* If the height of the edge is not a multiple of the threshold, $V/C = n + f$ with $0 < f < 1$, pixel i emits $n+1$ events when $a_i \le fC$ and n otherwise. Reproducing the count exactly costs the binary entropy of f per pixel, and a zero-rate statement then holds only for a distortion measure that tolerates one unmatched event per pixel, with a modified threshold. *Sensor.* Noise events, threshold mismatch between pixels, a refractory time and a jitter that depends on the slope are outside the model, and their effect on the three regimes is not analyzed here.

V. Moved Events

The decoder (22) needs the edge and the flow. A sender has neither. It has the events, and the closest it can come to (22) is to let the receiver move events it already holds along the motion. This is motion-compensated prediction, the core of a hybrid video coder [22], [23]: a block of a decoded picture, displaced by a motion vector, predicts a block of the next picture, and the coder sends the residual. This section gives what a moved event costs in timing error and in the fidelity of a receiver that counts events.

A. Timing Error

Proposition 2 (Moved events)**.** *Let the same edge pass pixel* j *in an earlier window and a different pixel* i *later, under (A1) to (A4) and (A2′), and let the motion be exact, so that the* k*th event of pixel* j *is moved to pixel* i *at the time* $\tilde{T}_{i,k} = T_{j,k} + t_i - t_j$*. Then*

$$\frac{1}{n}\,\mathbb{E}\|T_i - \tilde{T}_i\|^2 = D_{A,i} + D_{A,j} + \frac{4n^2 - 1}{12}\,(\tau_i - \tau_j)^2. \tag{28}$$

With equal event spacings the error is $2D_A$, twice that of (22): a moved event carries the phase and the jitter of its source pixel, and both are independent of those of its target. Unequal spacings, as under a change of scale, add the last term, which grows with the square of the number of events per pixel. If the moved events fall on their own pixel, as for an object at rest, the phases coincide and only the jitter term $2\sigma^2$ remains.

B. Voxel Fidelity

The timing error (14) presupposes that predicted and true events are paired. Receivers of event streams count events instead, on a voxel grid: the number of events of each polarity in blocks of $\Delta_x \times \Delta_x$ pixels and in time bins of length Δ_t. A video coder codes the same counts as pictures. For a reconstruction with voxel counts c' and a true set with counts c, of sizes N' and N, the fidelity is

$$F = \frac{2\sum \min(c', c)}{N' + N} = 1 - \frac{\sum |c' - c|}{N' + N}, \tag{29}$$

the share of the events of the two sets that can be paired one to one inside a voxel.

For blocks of one pixel, a train that is long against the event spacing, and a jitter that is small against the bin, Lemma 2 gives the fidelity of the regenerated events (22) directly: a bin holds at most one event of each train when $\Delta_t \le \tau_i$, the offset between the two trains is uniform, and in expectation over the phase

$$F = \frac{\Delta_t}{\tau_i} \qquad \text{for } \Delta_t \le \tau_i, \tag{30}$$

with $F = 1$ at $\Delta_t = \tau_i$. The voxel counts of a pixel are thus reproduced only when the bin is at least as long as its event spacing, the same scaling as the threshold of Theorem 2, for which $\sqrt{D_{A,i}} = \tau_i/\sqrt{12}$ when $\sigma \ll \tau_i$.

Moved events do almost as well on this grid. With equal spacings and no jitter, the kth moved event lies at the offset $\delta = A_i - A_j$ from the kth true event, and δ has the triangular density $(\tau_i - |\delta|)/\tau_i^2$ on $(-\tau_i, \tau_i)$. Two events at the offset x share a bin with probability $[1 - |x|/\Delta_t]^+$ over the position of the bins. When $|\delta| > \tau_i - \Delta_t$, a moved event can share its bin with a neighbor of its true event instead, which has a partner for $n-1$ of the n events of a train. For a train of n events and $\Delta_t \le \tau_i/2$, the two contributions give

$$F = \frac{\Delta_t}{\tau_i} - \frac{\Delta_t^2}{3n\tau_i^2}. \tag{31}$$

A voxel grid does not distinguish the events of a pixel, so the doubled timing error of (28) lowers the fidelity only by a term of order $1/n$.

For prediction, (31) sets a limit that no motion estimate can raise. With the exact motion, the counts of the previous window moved to their targets reproduce a share of about Δ_t/τ_i of the true counts when the bin is shorter than the event spacing, and the rest is a residual. On such a bin the count of a pixel is zero or one, and which bins hold the ones is decided by the phases, which motion does not carry. When the bin is long against the spacing, a bin holds whole trains or long parts of them, the offset δ moves only the events at its ends, and the moved counts approach the true ones. A motion-compensated coder of voxel counts can therefore save bits through its prediction on long bins and not on short ones.

VI. Experiments

The experiments address three questions: on the model, the three regimes of Theorem 2 (Section VI-B); on the model and on recordings, whether moved counts follow (31) (Sections VI-C and VI-D); and what a standard video coder spends against these limits on recordings (Section VI-E).

A. Protocol

Data. The recordings from a static camera are the 112 recordings of the training split of eTraM [24], 61 by day and 51 by night, from a sensor of 1280×720 pixels above a road, with 2.4×10^{10} events. The recordings from a moving camera are the 41 training sequences of DSEC [25], left event camera, 640×480 pixels, on a car, with 3.7×10^{10} events in 2634 s. A window of 33.3 ms of the median DSEC sequence holds 1.15 events per pixel, against 0.04 for eTraM. Twelve eTraM recordings and five DSEC sequences are held out for tuning, and the results below are over the other 100 and 36.

Grids. Time is cut into windows of $T_0 = 33.3$ ms and each window into k slices, so $\Delta_t = T_0/k$. A voxel holds $s \times s$ pixels and one polarity. We use $s \in \{1, 2\}$ and $k \in \{1, 3, 6, 12, 32, 64\}$, from $\Delta_t = 33.3$ ms to 0.52 ms. A group is four consecutive windows, taken every 3 s of an eTraM recording and every 2 s of a DSEC sequence.

Video coder. The counts of a group form one clip of $4k$ pictures, with the two polarities side by side in the luma plane of a 4:2:0 picture and constant chroma. They are coded by x265 3.5 [27], an encoder of HEVC [22], through ffmpeg 4.4.2, in the Main profile at the preset `slow`. Each clip has one intra picture and P pictures after it, without B pictures. Adaptive quantization and the psychovisual options are off. In lossless mode the luma is the count. At a constant quantization parameter (QP) the luma is 16 times the count, the decoder rounds the luma divided by 16, and the intra picture is coded at the QP of the P pictures. Counts are clipped at 15. On one-pixel voxels with $k \geq 3$, clipping removes less than 1% of the events; at $k = 1$ and on two-pixel voxels it removes up to 6% and 28%, by night. Every bit of the HEVC stream is counted, parameter sets included.

Voxel list. As a coder without prediction, the voxel list sorts the events of a window by voxel, writes the gaps between their voxel indices in a variable-length integer code and compresses the result with Zstandard at level 19 [28]. It codes the same clipped counts as HEVC and returns them exactly.

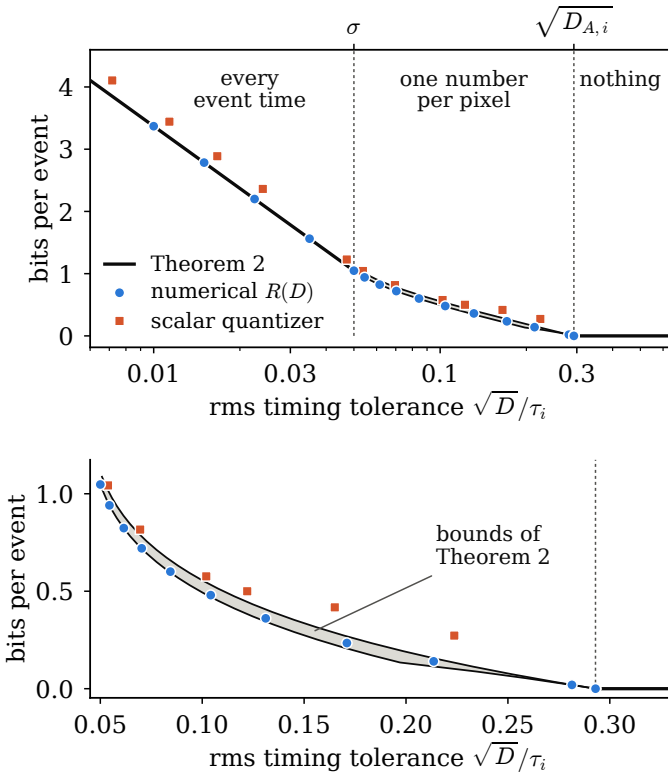


Fig. 2. Rate of one pixel against the timing tolerance for $n = 3$ and $\sigma/\tau_i = 0.05$: Theorem 2, a numerical rate-distortion function and a scalar quantizer. Bottom: the regime between σ and $\sqrt{D_{A,i}}$.

Tasks. With the annotated objects of eTraM and of DSEC-Detection [29], the fidelity (29) inside the boxes of the objects (box F), the share of boxes that the decode leaves without events, and the error of the motion of each box estimated by contrast maximization [30], in pixels per window.

Statistics. Bits are per true event. A recording enters with the sums over its groups, and we report the median over recordings. The coder comparisons of Section VI-E use one group in 40 of eTraM and one in 20 of DSEC, with all grids and QPs (42 recordings by day, 34 by night and 36 sequences); the representation of Section VI-E uses one group in five.

B. The Three Regimes

Fig. 2 compares Theorem 2 with a rate-distortion function that is computed without it, for one pixel with $n = 3$ and $\sigma/\tau_i = 0.05$. These values illustrate the theorem and are not those of a particular sensor. By (32), the event times are an orthogonal image of the scalar X_i, which carries the phase, and of $n - 1$ normal variables. The iteration of [31] gives the rate-distortion function of X_i on a grid of 1217 points, and the normal components enter in closed form at the same slope.

Below σ^2 the numerical rate equals (17) to within 4×10^{-6} nat. Between σ^2 and $D_{A,i}$ it stays inside the band between the bounds (18) and (21), which is at most $\Gamma_i = 0.187$ bit per pixel wide here, against the limit of 0.255 bit. The rate leaves

TABLE I
MOVED COUNTS AND LOSSLESS HEVC ON THE MODEL, $n = 3$, ONE-PIXEL VOXELS, MEANS OVER THREE SEEDS. F: THE COUNTS OF THE PREVIOUS WINDOW MOVED BY THE TRUE MOTION. b_{P}: BITS OF THE P PICTURES PER EVENT, X265 LOSSLESS.

τ (ms)	k	Δ_t/τ	F	(31)	b_{P}
1	1	33.3	0.990	–	0.14
1	3	11.1	0.970	–	0.45
1	12	2.78	0.879	–	1.95
1	32	1.04	0.851	–	5.39
1	64	0.521	0.489	–	7.89
4	12	0.694	0.639	–	3.76
4	32	0.260	0.251	0.253	7.01
4	64	0.130	0.127	0.128	9.48
8	12	0.347	0.331	0.334	5.06
8	64	0.065	0.064	0.065	10.13

σ^2 on the lower bound $U_i - \Gamma_i$: the phase is still described finely there, and the code profits fully from its distribution, which is uniform and not normal. Toward $D_{A,i}$ the description becomes coarse, the rate depends on the phase only through its variance, and the numerical rate, the bound (21) and U_i meet. From $D_{A,i}$ on the rate is zero, and the regenerated events (22) of 2×10^6 simulated pixels have the distortion $0.9994\, D_{A,i}$.

The squares are a uniform scalar quantizer with entropy coding, applied to what each regime describes. On every component of the event times it stays 0.254 to 0.258 bit per event above (17). This is the loss $\frac{1}{2}\ln(2\pi e/12)$ of a uniform quantizer at fine resolution, the difference between the entropy $h - \ln\Delta$ of the index of a quantizer of step Δ and (17) at the distortion $\Delta^2/12$. On the mean event time alone, with the other components left out, it stays 0.09 to 0.19 bit per event above the lower bound up to $D_{A,i}$. A code that describes every event time, then one number per pixel, then nothing, is therefore within 0.26 bit per event of the limit in all three regimes.

C. Moved Counts on the Model

Straight edges cross a sensor of 320×240 pixels at six pixels per window, 40 pixels apart, for twelve windows. Every row of pixels is a path of the model, with τ from 0.25 to 8 ms, $n \in \{1, 3\}$ and $\sigma = 0.05\,\tau$. The prediction of window m is the counts of window $m-1$ moved by exactly six pixels, the best that block motion compensation can do.

Wherever (31) applies, in all twelve cells of τ, n and k with $\Delta_t \le \tau/2$, the fidelity of the moved counts lies 0.001 to 0.005 below it (Table I). As Δ_t/τ falls from 33 to 0.52 at $\tau = 1$ ms, the fidelity falls from 0.990 to 0.489, so half of the counts are a residual even under the true motion.

x265 in lossless mode codes the same scenes and follows the same limit. Its P pictures cost about 0.1 bit per event at $\Delta_t/\tau \ge 30$, about 1 to 3 at $\Delta_t/\tau \approx 3$, and 5 to 14 at $\Delta_t/\tau \le 1$. They are coded with merge and inter prediction in about equal parts, and fewer than 9% of their coding units are skipped. The cost of a P picture thus rises where (31) says that the moved counts stop predicting the true ones.

D. Moved Counts on Recordings

On recordings neither the motion nor the event spacing is known. Each pixel estimates its spacing τ from the intervals between its events, the motion of each block of 20×20 pixels is chosen as the shift that best predicts window w from window $w-1$, and the same shift then predicts window $w+1$ from window w, so that the prediction is causal. A time offset per block takes up motion that is not a whole number of pixels per window. The fidelity of the prediction is summed in bins of Δ_t/τ and compared with a prediction that has the right pixels and random times, and with (31). Pixels whose spacing matches a flicker of 100 or 120 Hz are reported apart, and hot pixels are masked. *[To be completed: results of directive 020 and Fig. 3.]*

Figure file `figs/fig_threshold.pdf` (pending)

Fig. 3. Fidelity of the causally predicted counts against Δ_t/τ, with τ measured per pixel, for eTraM by day and by night and for DSEC; a prediction with the right pixels and random times; and (30) and (31).

TABLE II
LOSSLESS BITS PER EVENT, X265 (b_{L}) AND THE VOXEL LIST (b_{Q}), AND THE MEDIAN RATIO PER RECORDING.

	eTraM, day			eTraM, night			DSEC		
(s,k)	b_{L}	b_{Q}	L/Q	b_{L}	b_{Q}	L/Q	b_{L}	b_{Q}	L/Q
$(1,1)$	2.13	2.18	1.00	1.97	2.01	0.98	1.64	1.69	0.97
$(1,3)$	3.32	3.40	0.98	3.14	3.25	0.97	2.76	2.83	0.99
$(1,12)$	6.43	6.12	1.06	5.51	5.59	1.00	4.88	4.74	1.04
$(1,64)$	11.5	10.2	1.16	9.75	9.06	1.11	8.09	7.48	1.09
$(2,3)$	2.08	2.17	0.95	1.93	1.99	0.98	1.49	1.59	0.93
$(2,64)$	8.87	7.76	1.18	7.52	6.96	1.12	5.96	5.36	1.10

E. A Standard Video Coder Against the Limits

HEVC on count pictures has been used to code event streams without loss [11], and we use x265 here as a probe: it has motion-compensated prediction, a transform and a quantizer, and its own record of how it codes each picture shows where its bits go.

Without loss. x265 needs no fewer bits than the voxel list, which predicts nothing (Table II). Over the twelve grids, the ratio of the two is 0.92 to 1.02 at slices of $T_0/6$ or longer, 1.00 to 1.07 at $T_0/12$, and 1.06 to 1.18 at $T_0/32$ and $T_0/64$. From $k = 3$ on it rises with k on every stratum and at both voxel sizes. The intra picture takes 26% to 31% of the lossless bits at $k = 1$ and less than 1.5% at $k = 64$, so the P pictures carry the cost on fine grids. They do use prediction: merge accounts for 66% to 95% of their coding units on eTraM, and inter prediction for 68% to 87% on DSEC, where the whole scene moves. The prediction does not reduce the bits below those of a list of the counts. This is what Theorem 1 and (31) lead one to expect: the motion is shared by the pixels of a path, and a lossless code of the counts must still describe the phases, which decide the counts on fine grids and which the motion does not carry.

With loss. The quantizer of a video coder acts on the amplitude of the luma. On a fine grid the counts are almost

all zero or one, so a coarser QP does not make a voxel less precise; it removes it. On the model at $\tau = 1$ ms and $k = 64$, the voxel list codes the counts exactly with 5.04 bits per event and lossless x265 with 7.93. x265 reaches $F = 0.964$ at QP 16 with 20.9 bits per event, 0.557 at QP 20 with 7.55, and 0.107 at QP 24 with 1.27: below the bits of the exact list, its fidelity is at most 0.27, since a voxel with one event has the luma 16, which vanishes between QP 14 and 26. At $k = 3$ on the same scenes, where a voxel holds up to a whole train, it reaches 0.982 with 0.60 bit per event and 0.959 with 0.15. On recordings the same holds. On the grid of two pixels and $T_0/3$, x265 by day reaches $F = 0.975$ with 3.26 bits per event at QP 14, 0.933 with 1.09 at QP 28, 0.840 with 0.31 at QP 36 and 0.706 with 0.051 at QP 46, with the intra picture at the default QP offset of x265. On fine grids the cheapest QP of the ladder, 28, still costs 0.2 to 1 bit per event, and at one bit per event by day the best point at $k = 64$ keeps a box fidelity of 0.011 and leaves 71% of the boxes without events, against 0.279 and 38% at $k = 12$.

Representation. The grid itself loses information before any coding. Turning the events into counts and back, with the events of a voxel placed at its first instant, loses nothing for the box fidelity at one pixel and at least three slices per window, with no box left empty. Two-pixel voxels raise the motion error: at $k = 3$ it is 4.01, 1.17 and 2.67 pixels per window by day, by night and on DSEC, against 1.49, 0.67 and 0.94 at one pixel.

Implications. The three measurements separate what a coder of event counts can gain from what it cannot. Prediction from motion pays on bins long against the event spacing and not on short ones, where the bits describe the phases and the jitter of Theorem 1. A quantizer of amplitudes degrades gracefully only where voxels hold several events, and on fine grids it removes single events instead. A receiver that needs fine timing therefore calls for a different description of the residual, by event times relative to a prediction rather than by amplitudes, and a receiver that is satisfied with a bin longer than the event spacing is served by motion compensation and the quantizer of a video coder as they are.

VII. Conclusion

The pixels on the path of a moving edge share the edge and its motion, and each has a threshold phase and a jitter of its own, which a lossless code must describe at every pixel. A timing tolerance removes first the jitter and then the phase from what a code must describe, and beyond a threshold set by the event spacing the events are regenerated from the edge and the motion alone. A sender that has only the events can move them along the motion, which doubles the timing error and, on a voxel grid, reproduces a share Δ_t/τ of the counts when the bin is shorter than the event spacing. A standard HEVC coder on count pictures behaves as these limits say: without loss its motion-compensated prediction saves nothing against a list of the counts, its P pictures grow expensive as the bin falls below the event spacing, and with loss its quantizer removes single events on fine grids. The theory covers one ramp edge in uniform passage, so texture, noise events and threshold mismatch remain to be modeled, and a coder that describes the residual of fine bins by event times remains to be built.

Appendix A
Proofs of Lemmas 1 and 2

Proof of Lemma 1. The reference starts in $a_i + C\mathbb{Z}$ by the definition of the phase. By (A1) it changes by $\pm C$, so it stays in this set, and an event requires $\ell_i = \mu_i \pm C$. □

Proof of Lemma 2. By (A2) and (A2′), ℓ_i is nondecreasing, equal to 0 before t_i, to $\gamma r_i(t - t_i)$ during the passage, and to V after it. Since $\mu_i \le 0 = \ell_i$ before the edge arrives and $\mu_i = a_i - C$, the events occur when ℓ_i reaches the levels $a_i + (k-1)C$, $k = 1, 2, \dots$, as long as the level is at most $V = nC$. Because $0 < a_i \le C$, exactly the levels $k = 1, \dots, n$ are reached. Solving $\gamma r_i(t - t_i) = a_i + (k-1)C$ gives (10), and a_i uniform on $(0, C]$ makes A_i uniform on $(0, \tau_i]$. □

Appendix B
Proof of Theorem 1

By Lemma 2, $T_i = c_i + A_i\mathbf{1} + \varepsilon_i$ with c_i a function of the template and the flow, $\mathbf{1}$ the all-ones vector, and $\varepsilon_i \sim \mathcal{N}(0, \sigma^2 I_n)$. Independence over i follows from (A3) and (A4). Let Q be an orthogonal matrix whose first row is $\mathbf{1}^\top/\sqrt{n}$. Then $Q(T_i - c_i) = (X_i, W_i)$ with

$$X_i = \sqrt{n}\, A_i + \varepsilon_i^{\parallel}, \qquad W_i \sim \mathcal{N}(0, \sigma^2 I_{n-1}), \tag{32}$$

where $\varepsilon_i^{\parallel} \sim \mathcal{N}(0, \sigma^2)$. The variables X_i and W_i are independent, because $Q\varepsilon_i$ has independent coordinates and A_i is independent of ε_i. Differential entropy is invariant under translations and orthogonal maps, so $h(T_i) = h(X_i) + h(W_i)$, which is (11).

Let f be the density of T_i and f_δ its average over the cells of the clock lattice. A cell has probability $\delta^n f_\delta$, so $H(Y_i \mid g, \phi) = n\ln(1/\delta) - \int f \ln f_\delta$ and $e_\delta = \int f \ln(f/f_\delta)$ is a relative entropy, which is nonnegative. The density f is a mixture over A_i of $\mathcal{N}(m, \sigma^2 I_n)$ densities and f_δ is the same mixture of their cell averages. By the joint convexity of relative entropy, e_δ is at most its largest value for a single $\mathcal{N}(m, \sigma^2 I_n)$ density, which is n times the value for a scalar normal density and depends only on δ/σ and on the offset of m in a cell. It tends to zero with δ/σ, uniformly in the offset [26, Ch. 8].■

Appendix C
Proof of Theorem 2

The distortion (14) is invariant under the translation by c_i and the orthogonal map Q, so R_i is the rate-distortion function of the independent components (X_i, W_i) of (32) under the total squared error nD. Their variances are s_i^2 and σ^2, the latter $n - 1$ times. We drop the index i and write η for $\eta(\sqrt{n}\,\tau_i, \sigma)$ and $\ln^+ x = [\ln x]^+$.

Lower bound (17) *and* (18). Let $(\hat{X}, \hat{W})$ have component errors d_X and $d_1, \dots, d_{n-1}$ with $d_X + \sum_j d_j \le nD$. Independence of the components gives $I(X, W; \hat{X}, \hat{W}) \ge I(X; \hat{X}) + \sum_j I(W_j; \hat{W}_j)$. For a Gaussian component, $I(W_j; \hat{W}_j) \ge \frac{1}{2}\ln^+(\sigma^2/d_j)$, and $I(X; \hat{X}) \ge h(X) - h(X - \hat{X}) \ge \eta -$

$\frac{1}{2}\ln(2\pi e\, d_X)$. The function $\ln^+(\sigma^2/d)$ is convex in d, so the Gaussian terms are smallest when all d_j equal a common d_W, and $d_W > \sigma^2$ brings no benefit. For $n \geq 2$ this gives

$$R(D) \;\geq\; \min_{d_W} \Big[\eta - \tfrac{1}{2}\ln\big(2\pi e\,(nD-(n-1)d_W)\big) + \tfrac{n-1}{2}\ln\frac{\sigma^2}{d_W}\Big], \tag{33}$$

with the minimum over $0 < d_W \leq \sigma^2$, $d_W < nD/(n-1)$. The bracket is convex in d_W and its derivative vanishes only at $d_W = D$. The minimum is at $d_W = D$ when $D \leq \sigma^2$, where it equals the right side of (17), and at $d_W = \sigma^2$ otherwise, where it equals $U(D) - \Gamma$. For $n = 1$ there is no W and the same values follow with $d_X = D$.

Equality for $D \leq \sigma^2$. Write $\varepsilon = \varepsilon' + Z$ with independent $\varepsilon' \sim \mathcal{N}(0, (\sigma^2 - D)I_n)$ and $Z \sim \mathcal{N}(0, DI_n)$, and let $\hat{T} = c + A\mathbf{1} + \varepsilon'$. Then $T = \hat{T} + Z$ with Z independent of $\hat{T}$, the distortion is D, and $I(T;\hat{T}) = h(T) - h(Z)$, which is the lower bound.

Upper bound for $\sigma^2 < D < D_A$. Spend no rate on W, at error σ^2 per component, and describe X at error $d = nD - (n-1)\sigma^2 \in (\sigma^2, s^2)$ through $\hat{X} = (1 - d/s^2)(X - \mathbb{E}X + Z') + \mathbb{E}X$, with Z' an independent normal variable of variance $s^2 d/(s^2 - d)$. The error is d, and $I(X;\hat{X}) = I(X; X + Z') \leq \frac{1}{2}\ln(s^2/d)$, because a normal input maximizes the mutual information across an additive normal noise channel at a given input variance [26].

Lower bound (21) *and zero rate.* For a variable Ξ with $\ln \mathbb{E}\, e^{\lambda(\Xi - \mathbb{E}\Xi)} \leq \lambda^2\varsigma^2/2$ for all λ under a law P, and any law Q, the nonnegativity of the relative entropy between Q and the law with density proportional to $e^{\lambda\Xi}$ with respect to P gives $D(Q\|P) \geq \lambda(\mathbb{E}_Q\Xi - \mathbb{E}_P\Xi) - \lambda^2\varsigma^2/2$, and the best λ gives $(\mathbb{E}_Q\Xi - \mathbb{E}_P\Xi)^2 \leq 2\varsigma^2 D(Q\|P)$. With P the law of X and Q its conditional law given $\hat{T}$, averaging over $\hat{T}$ yields $\operatorname{Var}\mathbb{E}[X \mid \hat{T}] \leq 2\varsigma^2 I(X;\hat{T})$. The uniform part of X has width $a = \sqrt{n}\,\tau$ and satisfies the condition with $a^2/12$, since $\ln(\sinh x/x) \leq x^2/6$, so X satisfies it with $\varsigma^2 = s^2$, and each W_j with $\sigma^2 \leq s^2$. Since the conditional mean minimizes the squared error and the components are independent,

$$\begin{aligned} nD &\geq \mathbb{E}\|T - \mathbb{E}[T \mid \hat{T}]\|^2 \\ &= nD_A - \operatorname{Var}\mathbb{E}[X \mid \hat{T}] - \sum_j \operatorname{Var}\mathbb{E}[W_j \mid \hat{T}] \\ &\geq nD_A - 2s^2 I(T;\hat{T}), \end{aligned} \tag{34}$$

which is (21). With no description the best output is the mean, which is (22), at distortion $\frac{1}{n}(s^2 + (n-1)\sigma^2) = D_A$, so (19) holds. ∎

Appendix D
Proof of Theorem 3

The vectors T_i are independent and the distortion is additive over i. For any $\hat{T}$ with average distortion at most D, let $D_i = \mathbb{E}\, d(T_i, \hat{T}_i)$. Then $I\big((T_i)_i;(\hat{T}_i)_i\big) \geq \sum_i I(T_i;\hat{T}_i) \geq \sum_i R_i(D_i)$ with $\frac{1}{M}\sum_i D_i \leq D$. Conversely, independent test channels whose rates are within any given margin of $R_i(D_i)$ come within M times that margin of the sum. Each R_i is convex, nonincreasing and continuous, and tends to infinity as $D_i \to 0$ by (17), so the minimum in (23) exists.

For $D < \sigma^2$, every R_i is differentiable at D with the derivative $-n/(2D)$ by (17), and it is convex on $(0,\infty)$. Hence $R_i(D_i) \geq R_i(D) - \frac{n}{2D}(D_i - D)$ for every $D_i > 0$, and summing over i with $\sum_i (D_i - D) \leq 0$ shows that the allocation $D_i = D$ is optimal. This gives (24) for $D < \sigma^2$, and the case $D = \sigma^2$ follows by continuity.

By Theorem 2, $\sum_i R_i(D_i) = 0$ if and only if $D_i \geq D_{A,i}$ for every i, and such an allocation with average at most D exists if and only if $D \geq \frac{1}{M}\sum_i D_{A,i} = \bar{D}_A$. Finally, (21) and $s_i^2 \leq n\tau_{\max}^2/12 + \sigma^2$ give $\sum_i R_i(D_i) \geq n\sum_i [D_{A,i} - D_i]^+/(2(n\tau_{\max}^2/12 + \sigma^2))$, and $\sum_i [D_{A,i} - D_i]^+ \geq M(\bar{D}_A - D)$, which is (26). ∎

Appendix E
Proof of Proposition 1

Divide side j of the box into the smallest number of equal cells of length at most $\varsigma_j = \sqrt{\xi}\,\sigma/(q\, S_j M^{\beta_j})$, and let $\hat{\theta}$ be the center of the cell that contains θ. Side j has at most $L_j M^{\alpha_j}/\varsigma_j + 1$ cells, so the index of $\hat{\theta}$ takes $\sum_j \ln\big(L_j S_j q\, M^{\alpha_j+\beta_j}/(\sqrt{\xi}\,\sigma) + 1\big)$ nats, which is (27). The decoder outputs $\hat{T}_{i,k} = \bar{t}_{i,k}(\hat{\theta})$. The segment from θ to $\hat{\theta}$ lies in the box, so by the mean value theorem $e_{i,k} = \bar{t}_{i,k}(\theta) - \bar{t}_{i,k}(\hat{\theta})$ satisfies $|e_{i,k}| \leq \sum_j S_j M^{\beta_j}\varsigma_j/2 \leq \sqrt{\xi}\,\sigma/2$. By (10), $T_{i,k} - \bar{t}_{i,k}(\theta) = A_i - \tau_i/2 + \varepsilon_{i,k}$ has mean zero given θ and variance $D_{A,i}$, and $e_{i,k}$ is a function of θ. Hence $\mathbb{E}\,(T_{i,k} - \hat{T}_{i,k})^2 = D_{A,i} + \mathbb{E}\, e_{i,k}^2 \leq D_{A,i} + \xi\sigma^2/4$, and the average over the path is at most $\bar{D}_A + \xi\sigma^2/4 \leq (1+\xi)\bar{D}_A$. ∎

Appendix F
Proof of Proposition 2

By (10) and (A4), $T_{i,k} - \tilde{T}_{i,k} = (A_i - A_j) + (k-1)(\tau_i - \tau_j) + \varepsilon_{i,k} - \varepsilon_{j,k}$. The phases and the jitter of the two pixels are independent, so this difference has the mean $(k - \frac{1}{2})(\tau_i - \tau_j)$ and the variance $\tau_i^2/12 + \tau_j^2/12 + 2\sigma^2 = D_{A,i} + D_{A,j}$. Its mean square is the variance plus the squared mean, and $\frac{1}{n}\sum_{k=1}^{n}(k - \frac{1}{2})^2 = (4n^2 - 1)/12$. ∎

References

[1] P. Lichtsteiner, C. Posch, and T. Delbruck. A 128×128 120 dB 15 μs latency asynchronous temporal contrast vision sensor. *IEEE Journal of Solid-State Circuits*, 43(2):566–576, 2008.

[2] G. Gallego, T. Delbruck, G. Orchard, C. Bartolozzi, B. Taba, A. Censi, S. Leutenegger, A. J. Davison, J. Conradt, K. Daniilidis, and D. Scaramuzza. Event-based vision: A survey. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 44(1):154–180, 2022.

[3] M. Miskowicz. Send-on-delta concept: An event-based data reporting strategy. *Sensors*, 6(1):49–63, 2006.

[4] J. Mark and T. Todd. A nonuniform sampling approach to data compression. *IEEE Transactions on Communications*, 29(1):24–32, 1981.

[5] P. Martínez-Nuevo, H.-Y. Lai, and A. V. Oppenheim. Delta-ramp encoder for amplitude sampling and its interpretation as time encoding. *IEEE Transactions on Signal Processing*, 67(10):2516–2527, 2019.

[6] A. Papoulis. Generalized sampling expansion. *IEEE Transactions on Circuits and Systems*, 24(11):652–654, 1977.

[7] A. A. Lazar and L. T. Toth. Perfect recovery and sensitivity analysis of time encoded bandlimited signals. *IEEE Transactions on Circuits and Systems I: Regular Papers*, 51(10):2060–2073, 2004.

[8] K. Adam, A. Scholefield, and M. Vetterli. Sampling and reconstruction of bandlimited signals with multi-channel time encoding. *IEEE Transactions on Signal Processing*, 68:1105–1119, 2020.

[9] K. Adam, A. Scholefield, and M. Vetterli. Asynchrony increases efficiency: Time encoding of videos and low-rank signals. *IEEE Transactions on Signal Processing*, 70:105–116, 2022.

[10] Z. Bi, S. Dong, Y. Tian, and T. Huang. Spike coding for dynamic vision sensors. In *2018 Data Compression Conference (DCC)*, pages 117–126, 2018.

[11] N. Khan, K. Iqbal, and M. G. Martini. Time-aggregation-based lossless video encoding for neuromorphic vision sensor data. *IEEE Internet of Things Journal*, 8(1):596–609, 2021.

[12] I. Schiopu and R. C. Bilcu. Lossless compression of event camera frames. *IEEE Signal Processing Letters*, 29:1779–1783, 2022.

[13] S. Banerjee, Z. W. Wang, H. H. Chopp, O. Cossairt, and A. K. Katsaggelos. Lossy event compression based on image-derived quad trees and poisson disk sampling. In *2021 IEEE International Conference on Image Processing (ICIP)*, pages 2154–2158, 2021.

[14] B. Huang and T. Ebrahimi. Event data stream compression based on point cloud representation. In *2023 IEEE International Conference on Image Processing (ICIP)*, pages 3120–3124, 2023.

[15] Daniel C. Stumpp, Himanshu Akolkar, Alan D. George, and Ryad B. Benosman. Flow-based visual stream compression for event cameras. *IEEE Internet of Things Journal*, 11(24):40229–40243, 2024.

[16] C. Brites and J. Ascenso. Neuromorphic vision data coding: Classifying and reviewing the literature. *IEEE Access*, 13:14626–14657, 2025.

[17] Zahra Rezaee, Catarina Brites, and João Ascenso. Lossy event compression: From event stream distortion to task performance. arXiv preprint arXiv:2608.28429, 2026.

[18] I. Rubin. Information rates and data-compression schemes for Poisson processes. *IEEE Transactions on Information Theory*, 20(2):200–210, 1974.

[19] S. Verdú. The exponential distribution in information theory. *Problemy Peredachi Informatsii*, 32(1):100–111, 1996. English translation: Problems of Information Transmission, vol. 32, no. 1, pp. 86–95.

[20] A. Lapidoth, A. Malär, and L. Wang. Covering point patterns. *IEEE Transactions on Information Theory*, 61(9):4521–4533, 2015.

[21] G. Koliander, D. Schuhmacher, and F. Hlawatsch. Rate-distortion theory of finite point processes. *IEEE Transactions on Information Theory*, 64(8):5832–5861, 2018.

[22] G. J. Sullivan, J.-R. Ohm, W.-J. Han, and T. Wiegand. Overview of the high efficiency video coding (HEVC) standard. *IEEE Transactions on Circuits and Systems for Video Technology*, 22(12):1649–1668, 2012.

[23] G. J. Sullivan and T. Wiegand. Rate-distortion optimization for video compression. *IEEE Signal Processing Magazine*, 15(6):74–90, 1998.

[24] Aayush Atul Verma, Bharatesh Chakravarthi, Arpitsinh Vaghela, Hua Wei, and Yezhou Yang. eTraM: Event-based traffic monitoring dataset. In *Proc. IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 22637–22646, 2024.

[25] M. Gehrig, W. Aarents, D. Gehrig, and D. Scaramuzza. DSEC: A stereo event camera dataset for driving scenarios. *IEEE Robotics and Automation Letters*, 6(3):4947–4954, 2021.

[26] T. M. Cover and J. A. Thomas. *Elements of Information Theory*. John Wiley & Sons, 2 edition, 2006.

[27] MulticoreWare. x265 HEVC encoder. https://x265.readthedocs.io/. Version 3.5.

[28] Y. Collet and M. Kucherawy. Zstandard compression and the 'application/zstd' media type. RFC 8878, Internet Engineering Task Force, 2021.

[29] DSEC-Detection. https://dsec.ifi.uzh.ch/dsec-detection/. Object labels for DSEC, left event camera.

[30] G. Gallego, H. Rebecq, and D. Scaramuzza. A unifying contrast maximization framework for event cameras, with applications to motion, depth, and optical flow estimation. In *Proc. IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 3867–3876, 2018.

[31] R. E. Blahut. Computation of channel capacity and rate-distortion functions. *IEEE Transactions on Information Theory*, 18(4):460–473, 1972.